

Assignment

March 6, 2022

```
[3]: import pandas as pd
import numpy as np
```

```
[4]: dataset_amazon=pd.read_csv('Amazon - Movies and TV Ratings.csv')
```

```
[5]: dataset_amazon.head()
```

```
[5]:
```

	user_id	Movie1	Movie2	Movie3	Movie4	Movie5	Movie6	Movie7	\
0	A3R50BKS70M2IR	5.0	5.0	NaN	NaN	NaN	NaN	NaN	
1	AH3QC2PC1VTGP	NaN	NaN	2.0	NaN	NaN	NaN	NaN	
2	A3LKP6WPMP9UKX	NaN	NaN	NaN	5.0	NaN	NaN	NaN	
3	AVIY68KEPQ5ZD	NaN	NaN	NaN	5.0	NaN	NaN	NaN	
4	A1CV1WROP5KTTW	NaN	NaN	NaN	NaN	5.0	NaN	NaN	

	Movie8	Movie9	...	Movie197	Movie198	Movie199	Movie200	Movie201	\
0	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	...	NaN	NaN	NaN	NaN	NaN	

	Movie202	Movie203	Movie204	Movie205	Movie206
0	NaN	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN

[5 rows x 207 columns]

```
[6]: dataset_amazon.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4848 entries, 0 to 4847
Columns: 207 entries, user_id to Movie206
dtypes: float64(206), object(1)
memory usage: 7.7+ MB
```

```
[7]: movie_data=pd.melt(dataset_amazon, id_vars=["user_id"], var_name="Movie_id",
    ↳value_name="Ratings")
movie_data
```

```
[7]:
```

	user_id	Movie_id	Ratings
0	A3R50BKS70M2IR	Movie1	5.0
1	AH3QC2PC1VTGP	Movie1	NaN
2	A3LKP6WPMP9UKX	Movie1	NaN
3	AVIY68KEPQ5ZD	Movie1	NaN
4	A1CV1WROP5KTTW	Movie1	NaN
...
998683	A1IMQ9WMFYKWH5	Movie206	5.0
998684	A1KLIKPUF5E88I	Movie206	5.0
998685	A5HG6WFZL010D	Movie206	5.0
998686	A3UU690TWXCG1X	Movie206	5.0
998687	AI4J762YI6S06	Movie206	5.0

[998688 rows x 3 columns]

1 1. Which movies have maximum views/ratings

```
[8]: top5=movie_data.groupby('Movie_id').agg({"Ratings":'count'}).reset_index().
    ↳sort_values(by='Ratings',ascending=False)
top5.columns=['Movie_id', '# Ratings']
top5.head(10)
```

```
[8]:
```

	Movie_id	# Ratings
31	Movie127	2313
46	Movie140	578
67	Movie16	320
5	Movie103	272
128	Movie29	243
197	Movie91	128
198	Movie92	101
194	Movie89	83
65	Movie158	66
10	Movie108	54

Highest Ratings

```
[9]: movie_data.groupby('Movie_id').agg({"Ratings":'sum'}).reset_index().
    ↳sort_values(by='Ratings',ascending=False).head(10)
```

```
[9]:
```

	Movie_id	Ratings
31	Movie127	9511.0

46	Movie140	2794.0
67	Movie16	1446.0
5	Movie103	1241.0
128	Movie29	1168.0
197	Movie91	586.0
198	Movie92	482.0
194	Movie89	380.0
65	Movie158	318.0
10	Movie108	252.0

```
[10]: movie_data.groupby('Movie_id').agg({"Ratings": 'sum'}).reset_index().
      ↪sort_values(by='Ratings',ascending=False)
```

```
[10]:
```

	Movie_id	Ratings
31	Movie127	9511.0
46	Movie140	2794.0
67	Movie16	1446.0
5	Movie103	1241.0
128	Movie29	1168.0
..
61	Movie154	1.0
160	Movie58	1.0
50	Movie144	1.0
170	Movie67	1.0
163	Movie60	1.0

[206 rows x 2 columns]

2. What is the average rating for each movie? Define the top 5 movies with the maximum ratings.

```
[11]: top5_avg=movie_data.groupby('Movie_id').agg({"Ratings": 'mean'}).reset_index()
      top5_avg.columns=['Movie_id', 'Average Ratings']
      top5_avg.head(10)
```

```
[11]:
```

	Movie_id	Average Ratings
0	Movie1	5.0000
1	Movie10	5.0000
2	Movie100	4.0000
3	Movie101	5.0000
4	Movie102	4.0000
5	Movie103	4.5625
6	Movie104	4.5000
7	Movie105	5.0000
8	Movie106	5.0000

9 Movie107 4.0000

```
[12]: ds_movie=pd.merge(top5, top5_avg, on='Movie_id')
      ds_movie.head(10)
```

```
[12]:
```

	Movie_id	# Ratings	Average Ratings
0	Movie127	2313	4.111976
1	Movie140	578	4.833910
2	Movie16	320	4.518750
3	Movie103	272	4.562500
4	Movie29	243	4.806584
5	Movie91	128	4.578125
6	Movie92	101	4.772277
7	Movie89	83	4.578313
8	Movie158	66	4.818182
9	Movie108	54	4.666667

3 3. Define the top 5 movies with the least audience.

```
[13]: ds_movie.sort_values(['Average Ratings'], ascending=False).sort_values(['#_
      ↳Ratings'], ascending=True)
```

```
[13]:
```

	Movie_id	# Ratings	Average Ratings
205	Movie1	1	5.000000
200	Movie153	1	5.000000
203	Movie36	1	5.000000
199	Movie152	1	5.000000
202	Movie37	1	5.000000
..
4	Movie29	243	4.806584
3	Movie103	272	4.562500
2	Movie16	320	4.518750
1	Movie140	578	4.833910
0	Movie127	2313	4.111976

[206 rows x 3 columns]

Data Considerations - All the users have not watched all the movies and therefore, all movies are not rated. These missing values are represented by NA. - Ratings are on a scale of -1 to 10 where -1 is the least rating and 10 is the best. - Recommendation Model: Some of the movies hadn't been watched and therefore, are not rated by the users. Netflix would like to take this as an opportunity and build a machine learning recommendation algorithm which provides the ratings for each of the users.

```
[14]: # executed below in Command prompt
      # conda install -c conda-forge scikit-surprise
```

```
[15]: !pip install scikit-surprise
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: scikit-surprise in /usr/local/lib/python3.7/site-packages (1.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/site-packages (from scikit-surprise) (0.14.1)
Requirement already satisfied: numpy>=1.11.2 in /usr/local/lib/python3.7/site-packages (from scikit-surprise) (1.18.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.7/site-packages (from scikit-surprise) (1.4.1)
Requirement already satisfied: six>=1.10.0 in /usr/local/lib/python3.7/site-packages (from scikit-surprise) (1.14.0)
```

```
[16]: from surprise import Reader
      from surprise import Dataset
      from surprise import KNNBasic, SVD
      from surprise import accuracy
      from surprise.model_selection import train_test_split
```

```
[17]: movie_data.fillna(0, inplace=True)
```

```
[18]: reader = Reader(line_format='user item rating', rating_scale=(-1, 10))
```

```
[19]: Dataset=Dataset.load_from_df(movie_data, reader=reader)
```

```
[20]: train_set, test_set=train_test_split(Dataset, test_size=.25, random_state=10)
      ↪ #Split data into training and testing set
```

```
[21]: #KNN MODEL
      #Train the model
      sim_options = {'name': 'pearson', 'user_based': True}
      knn=KNNBasic(k=45, min_k=3, sim_options=sim_options)
```

```
[22]: knn.fit(train_set)
```

```
Computing the pearson similarity matrix...
Done computing similarity matrix.
```

```
[22]: <surprise.prediction_algorithms.knns.KNNBasic at 0x7f15c289ecd0>
```

```
[23]: predictions=knn.test(test_set)
```

```
[24]: result = pd.DataFrame(predictions, columns=['user_id', 'movie_id',
↳ 'base_ratings', 'predict_ratings', 'additional'])
result.drop('additional',axis=1,inplace=True)
```

```
[25]: result
```

```
[25]:
```

	user_id	movie_id	base_ratings	predict_ratings
0	A3VURT1JDRQQRI	Movie61	0.0	0.000000
1	A1FP2LA6M20JX0	Movie7	0.0	0.000000
2	A3BIIQC3A935LS	Movie46	0.0	0.022057
3	A39ZX6ML4X7E67	Movie34	0.0	0.000000
4	A2M5FI4CB6VUXF	Movie53	0.0	0.000000
...
249667	A1EBD2U23BP04Y	Movie128	0.0	0.000000
249668	A10A4PJBRIFGLE	Movie92	0.0	0.000000
249669	ACFK06NL9N7I8	Movie182	0.0	0.000000
249670	A1Y81TTIF0F5GX	Movie157	0.0	0.000000
249671	A1XYHEBZVXGKX1	Movie72	0.0	0.022057

[249672 rows x 4 columns]

```
[26]: accuracy.rmse(predictions)
```

RMSE: 0.3183

```
[26]: 0.31830539402296587
```

```
[27]: #SVD MODEL
#Train the model
svd=SVD()
```

```
[28]: svd.fit(train_set)
```

```
[28]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x7f156c82ead0>
```

```
[29]: predictions=svd.test(test_set)
```

```
[30]: result = pd.DataFrame(predictions, columns=['user_id', 'movie_id',
↳ 'base_ratings', 'predict_ratings', 'additional'])
result.drop('additional',axis=1,inplace=True)
```

```
[31]: result
```

```
[31]:
```

	user_id	movie_id	base_ratings	predict_ratings
0	A3VURT1JDRQQRI	Movie61	0.0	0.010079
1	A1FP2LA6M20JX0	Movie7	0.0	0.004391
2	A3BIIQC3A935LS	Movie46	0.0	0.003338

3	A39ZX6ML4X7E67	Movie34	0.0	0.008307
4	A2M5FI4CB6VUXF	Movie53	0.0	0.013513
...
249667	A1EBD2U23BP04Y	Movie128	0.0	0.004345
249668	A10A4PJBRIFGLE	Movie92	0.0	-0.086026
249669	ACFK06NL9N7I8	Movie182	0.0	-0.031928
249670	A1Y81TTIF0F5GX	Movie157	0.0	0.016468
249671	A1XYHEBZVXGKX1	Movie72	0.0	-0.012038

[249672 rows x 4 columns]

```
[32]: accuracy.rmse(predictions)
```

RMSE: 0.2794

```
[32]: 0.2794406865119178
```

```
[33]: #The best model is svd than knnbase due to lower RMSE value
```

```
[34]: def recommendation_movies(user):
    #get the list of unique movies
    movie_data=pd.melt(dataset_amazon, id_vars=["user_id"],
    var_name="Movie_id", value_name="Rating")
    movie_names=movie_data.Movie_id.unique()
    #movies watched by the user
    movies_watched=movie_data[movie_data["user_id"]==user].dropna().Movie_id
    #Movies the user didn't watch
    movies_not_watched=np.setdiff1d(movie_names, movies_watched)
    #Build the model
    model=svd.fit(Dataset.build_full_trainset())
    #Predictions
    movies_predict=[]
    for i in movies_not_watched:
        movies_predict.append((i, model.predict(user, i).est))
    return pd.DataFrame(movies_predict, columns=['Movie_id', 'Predictions']).
    sort_values(by='Predictions',ascending=False)['Movie_id'].head(10).to_list()
```

```
[35]: #Top 10 movies recommended for user=A3VURT1JDRQQRI
recommendation_movies('A3VURT1JDRQQRI')
```

```
[35]: ['Movie140',
'Movie205',
'Movie206',
'Movie182',
'Movie185',
'Movie158',
'Movie184',
```

```
'Movie196',  
'Movie204',  
'Movie173']
```

```
[36]: #Top 10 movies recommended for user=A1FP2LA6M20JX0  
recommendation_movies('A1FP2LA6M20JX0')
```

```
[36]: ['Movie140',  
       'Movie205',  
       'Movie206',  
       'Movie182',  
       'Movie185',  
       'Movie158',  
       'Movie184',  
       'Movie127',  
       'Movie196',  
       'Movie204']
```

```
[ ]:
```